Graphics for inference

- What is my model telling me?
- How can I tell other people?

Principles

- Graphs tell stories better than tables do
  - Use graphs to illustrate comparisons
  - Be careful about *units*
- Distinguish between (scientific) variables and (statistical) parameters
- Keep P values in their place
- What to do about raw data?

Smoking data

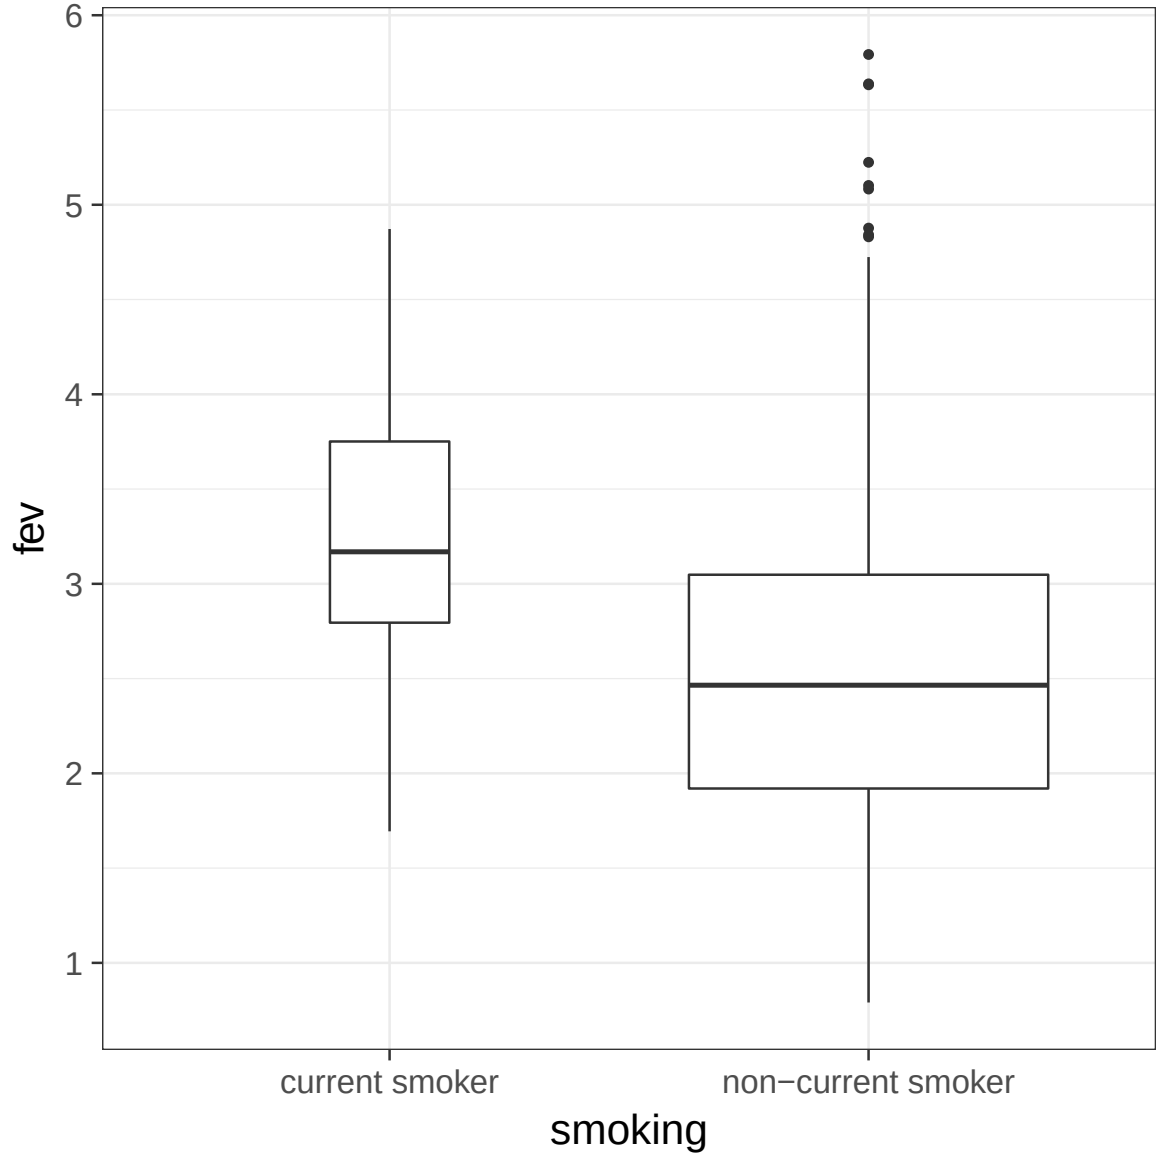

Smoking data

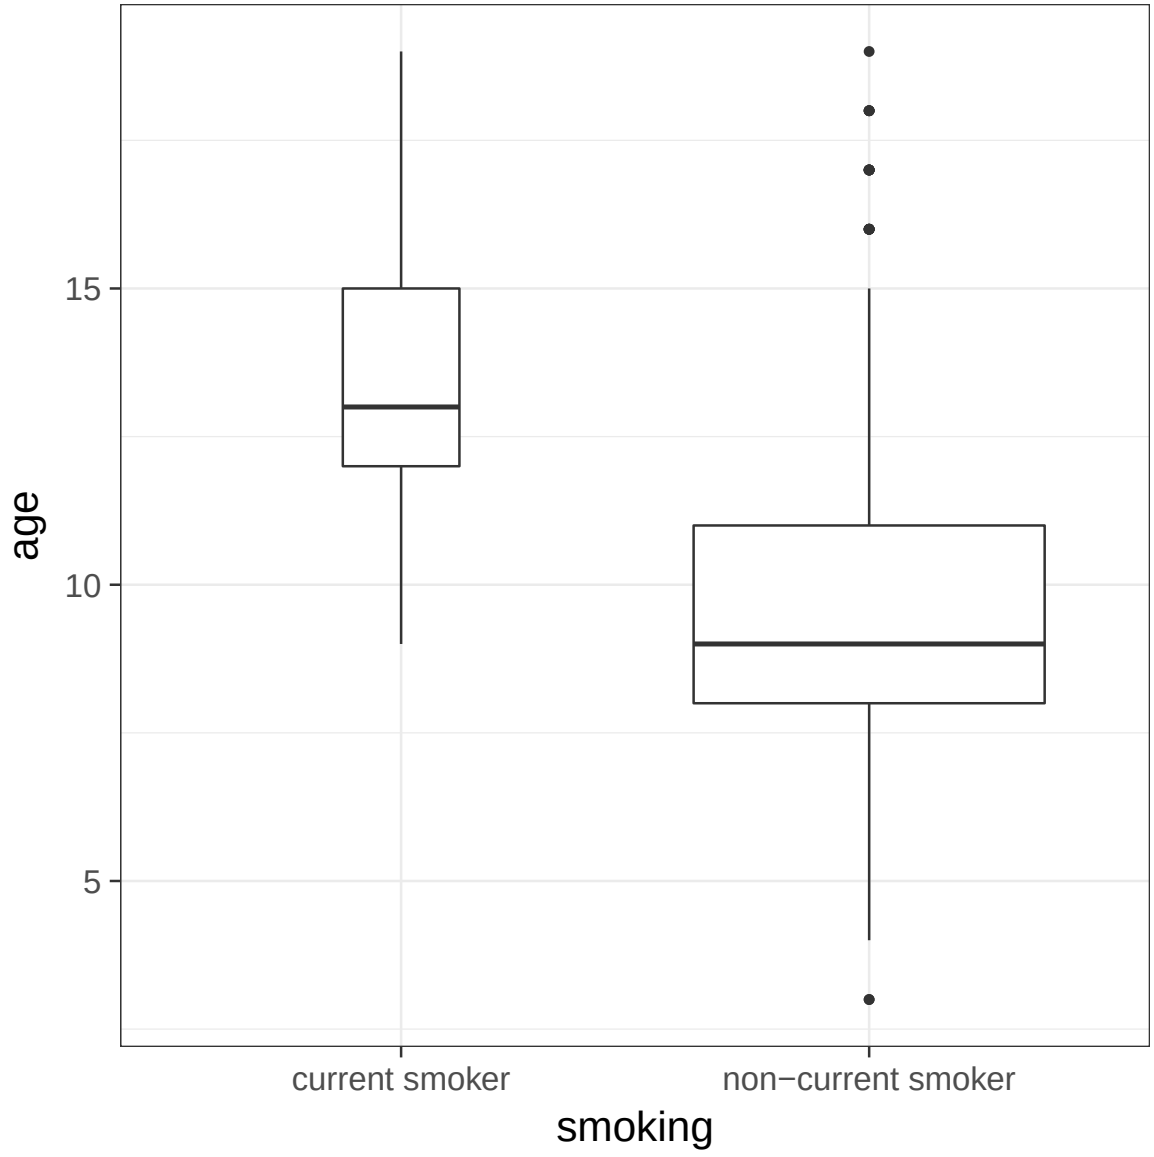

Smoking data

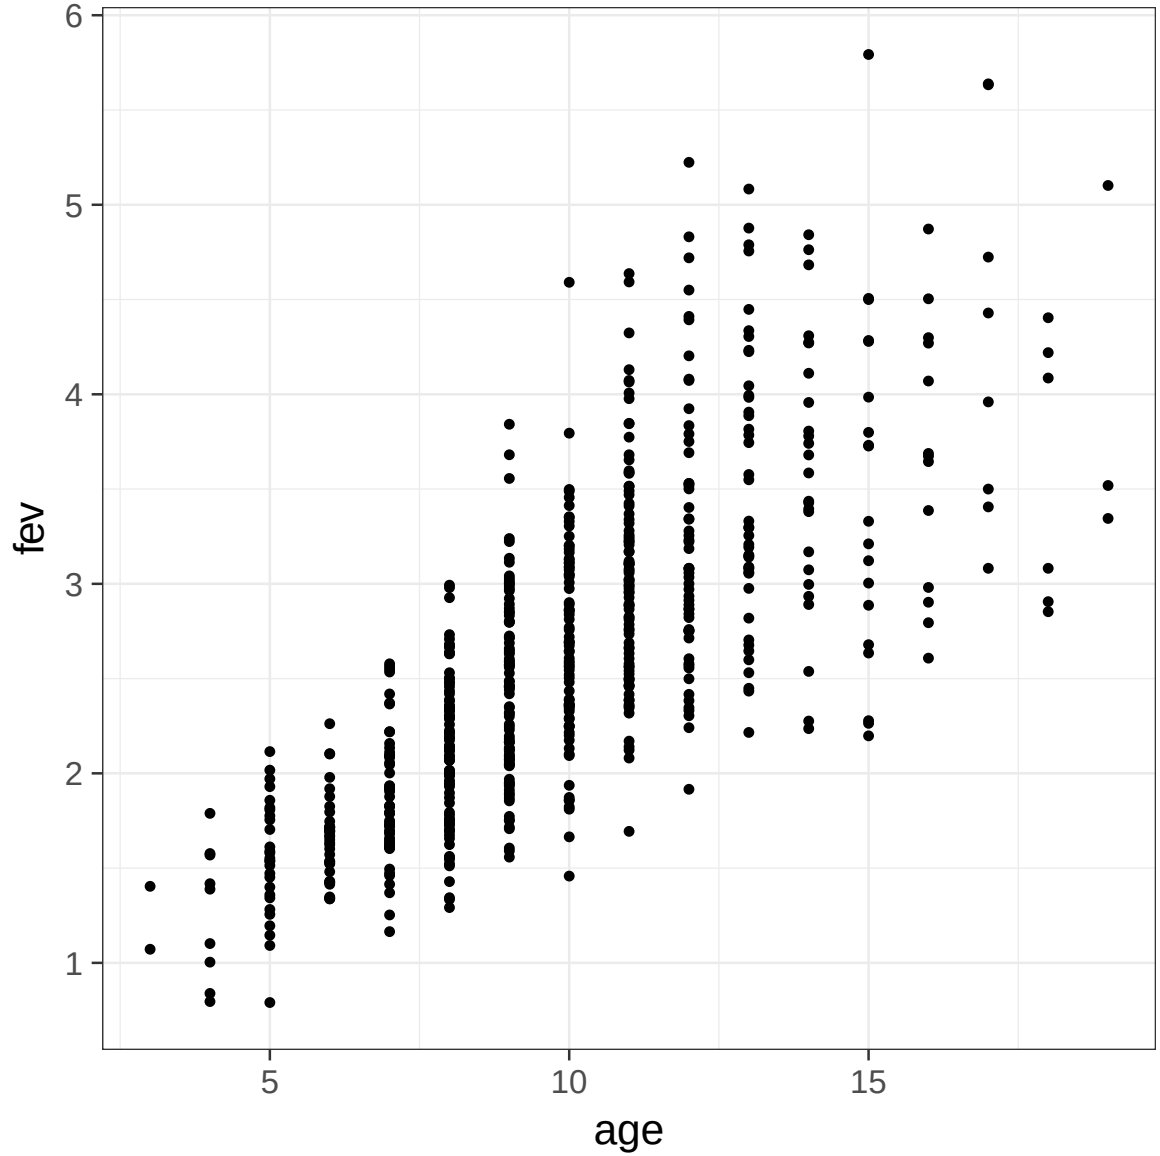

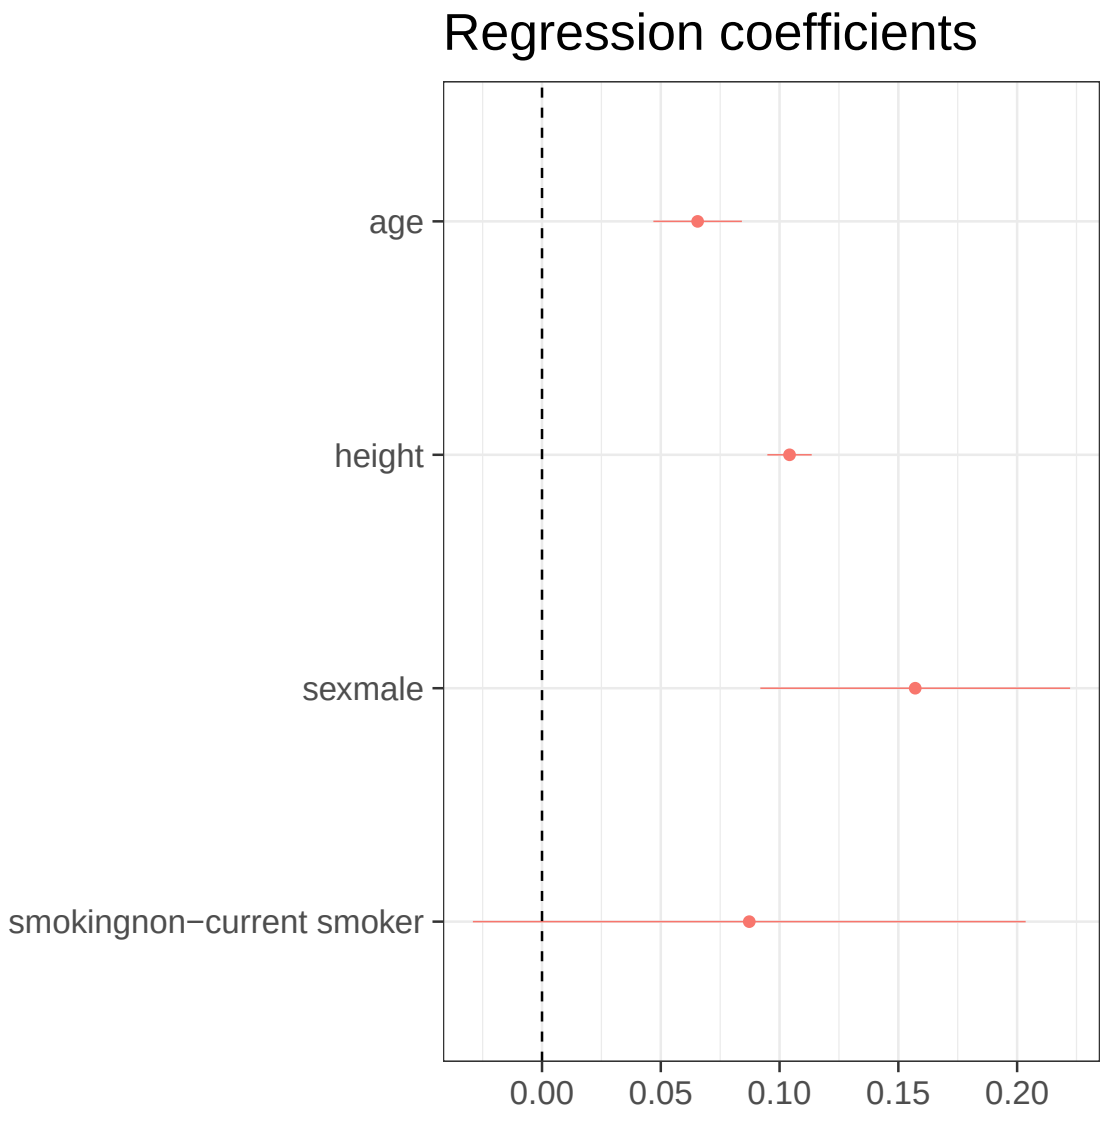
Regression coefficients
age
height
sexmale
smokingnon−current smoker
0.00
0.05
0.10
0.15
0.20

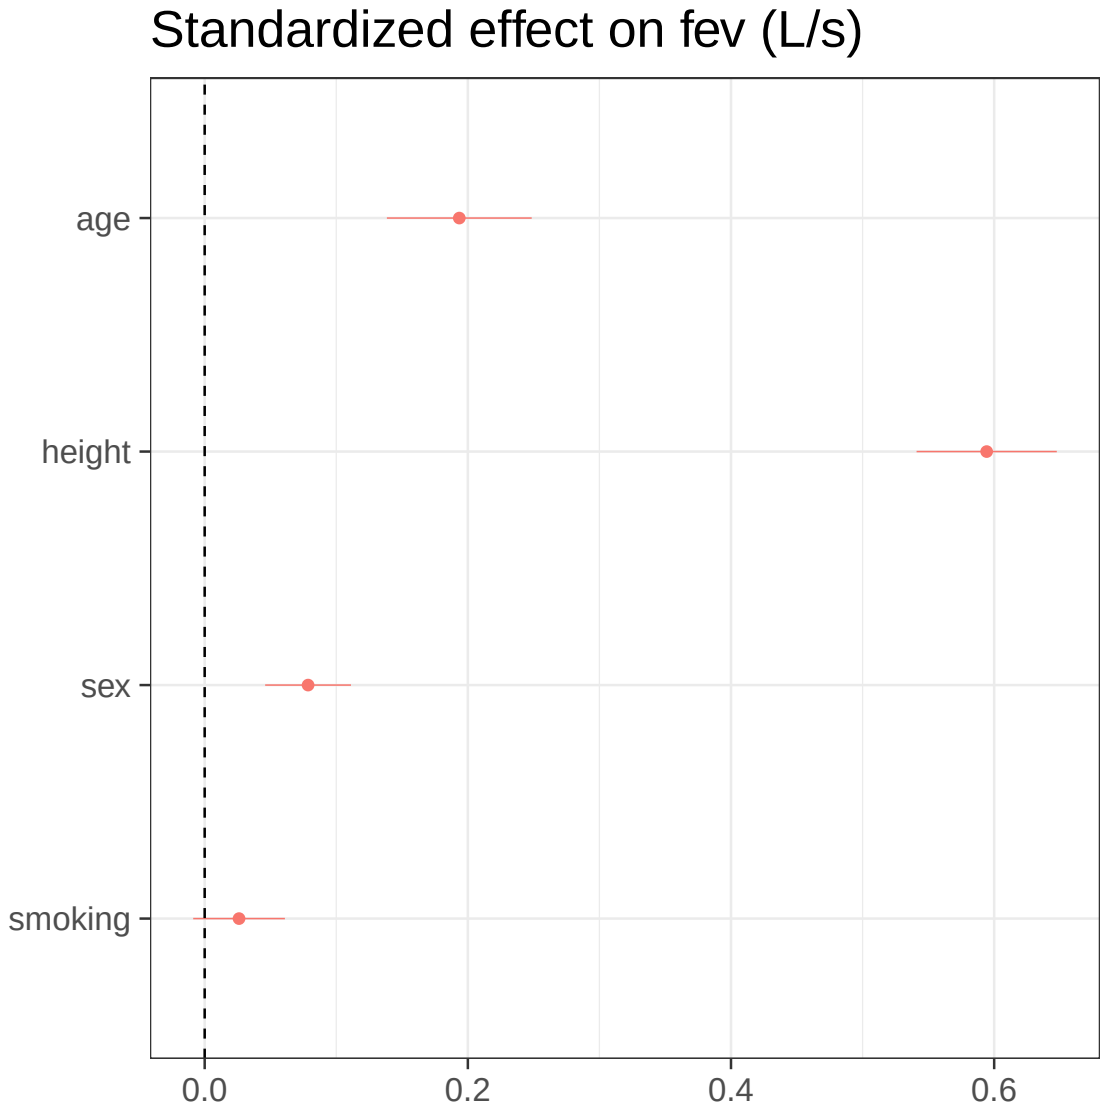
Standardized effect on fev (L/s)
age
height
sex
smoking
0.0
0.2
0.4
0.6

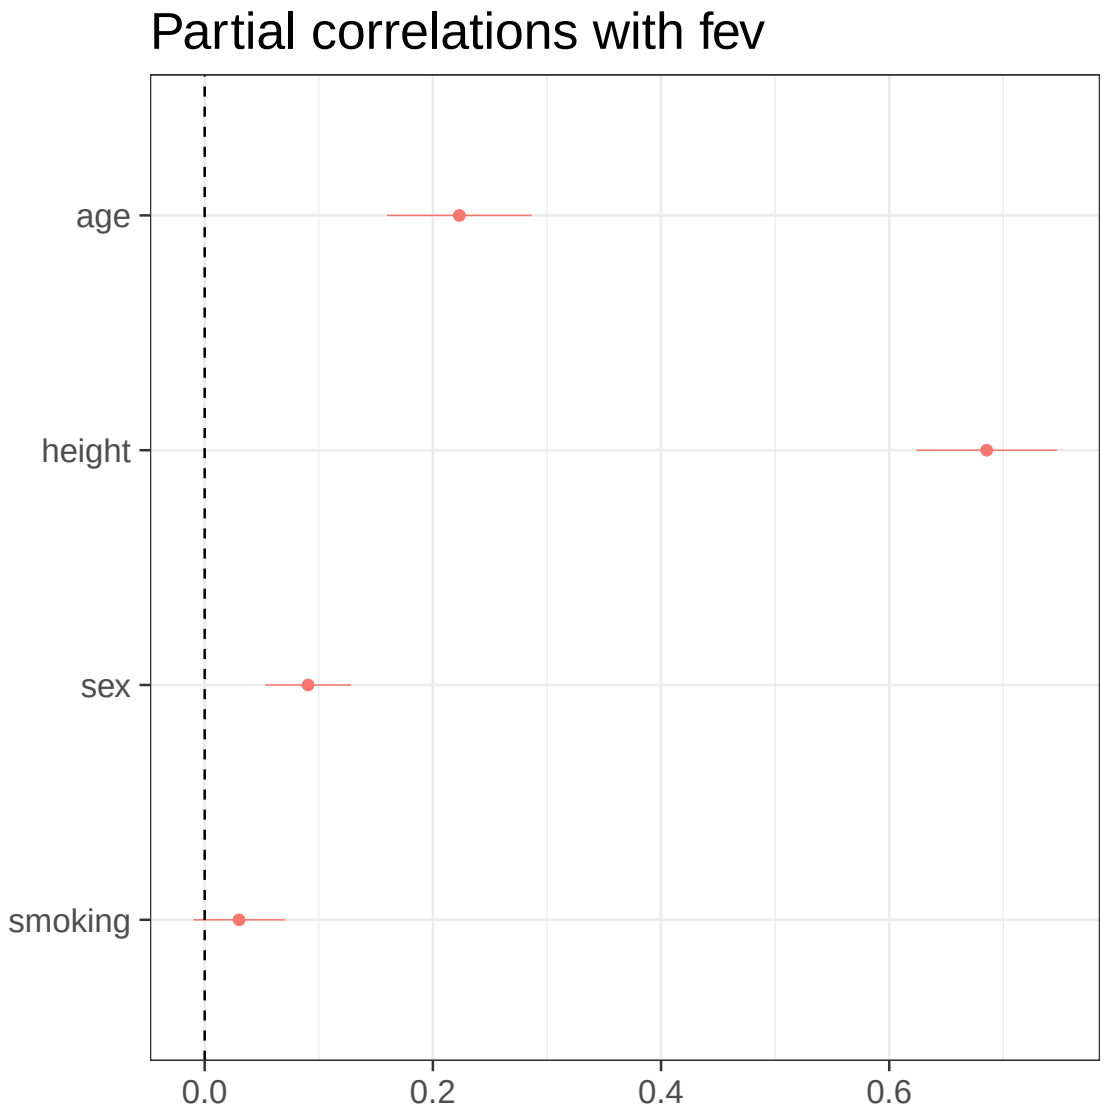
Partial correlations with fev
age
height
sex
smoking
0.0
0.2
0.4
0.6

Comparing effects on different response variables

- Put response variables on same scale:
  - Standardize
  - Logs
  - Proportions

Shape of response

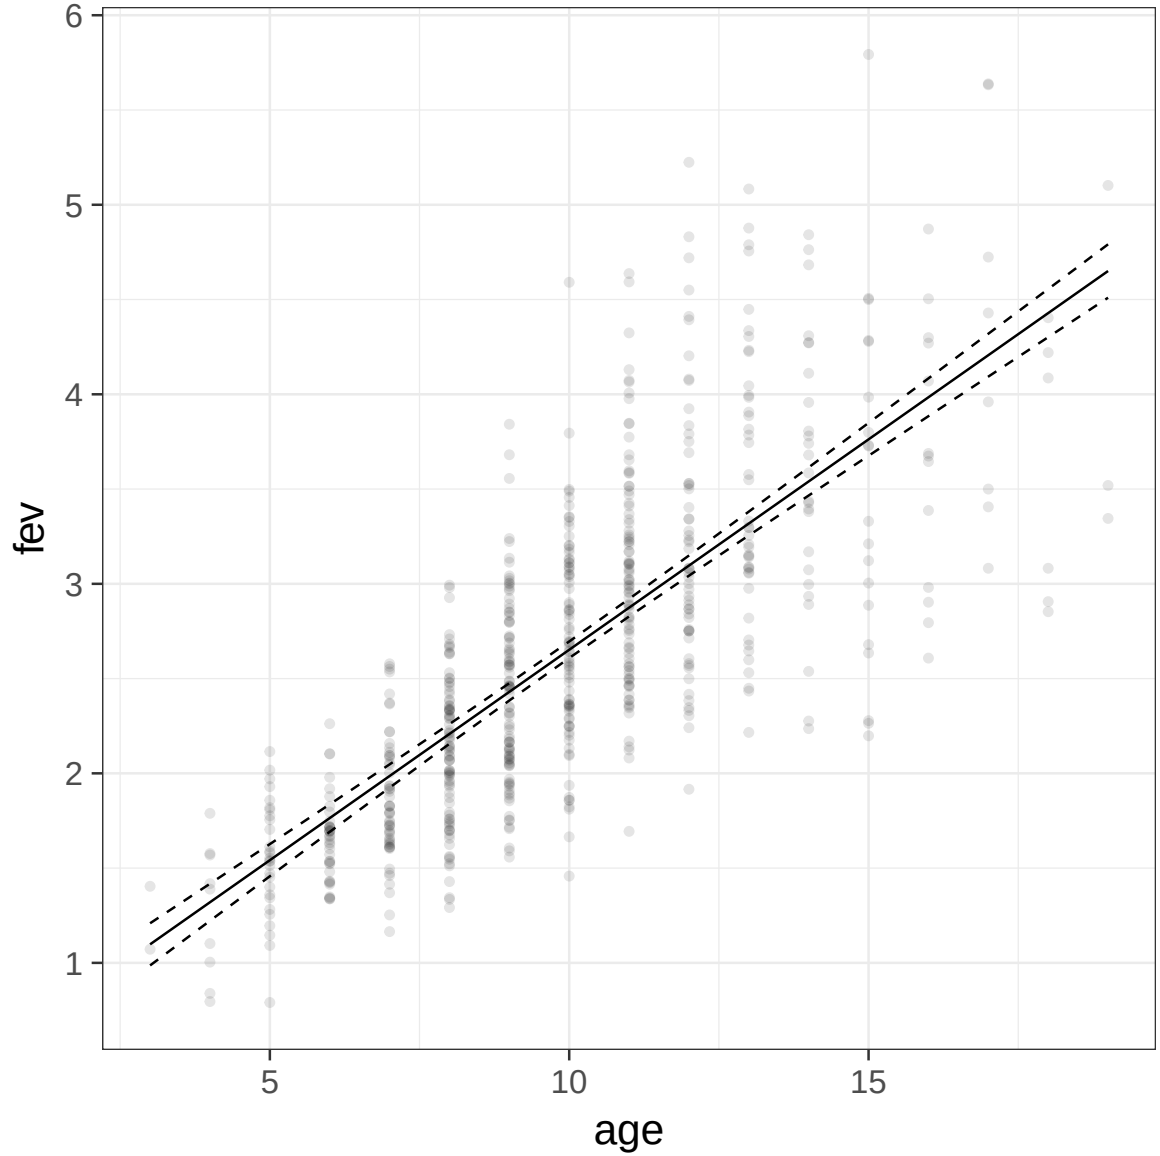

Standard prediction plot

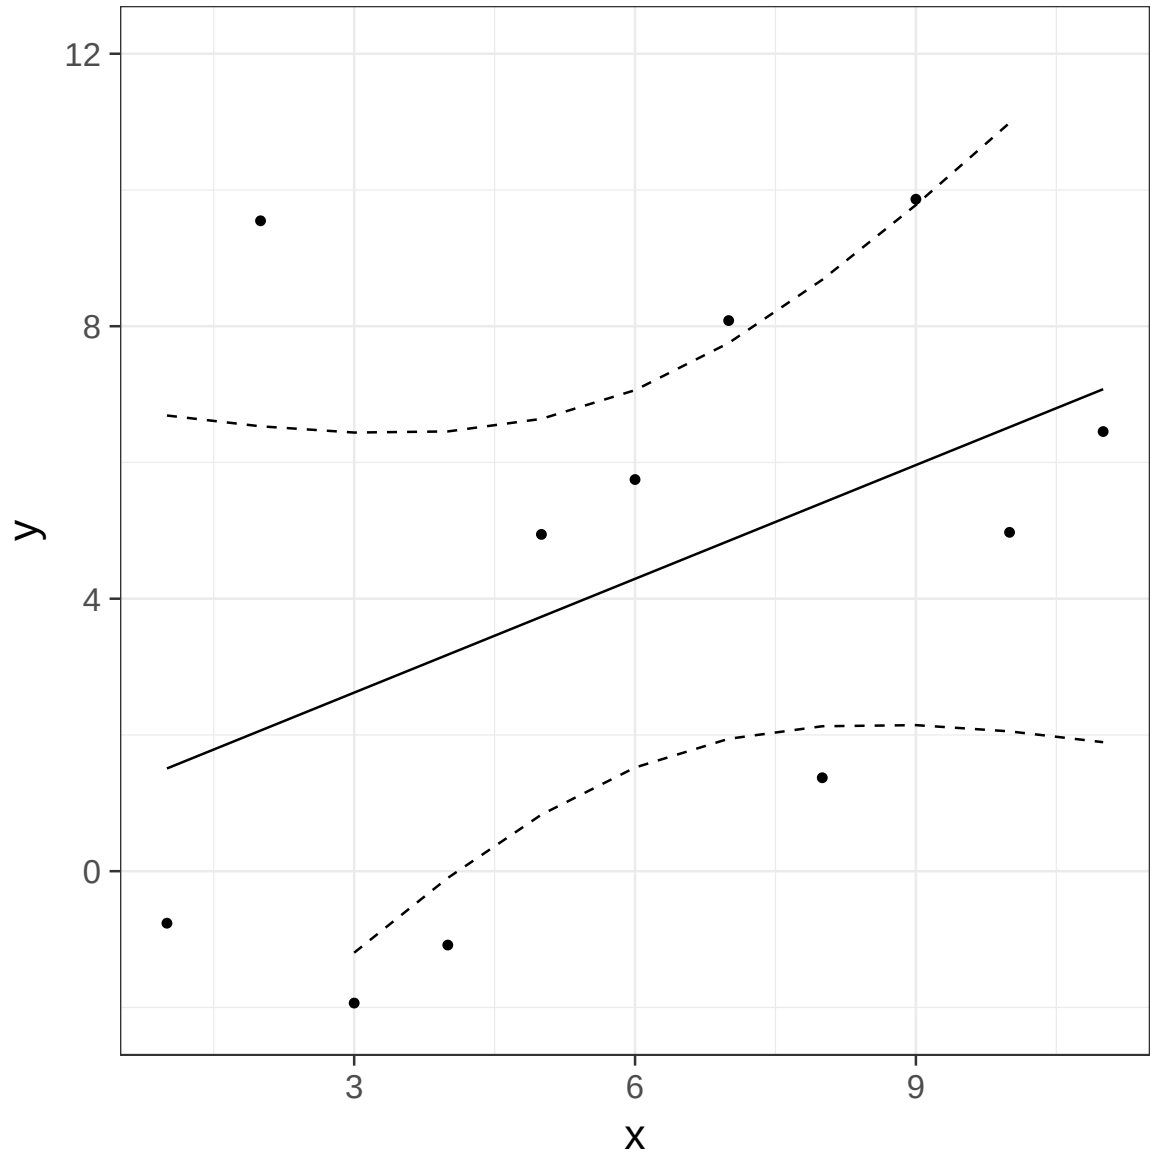

Marginal prediction plot

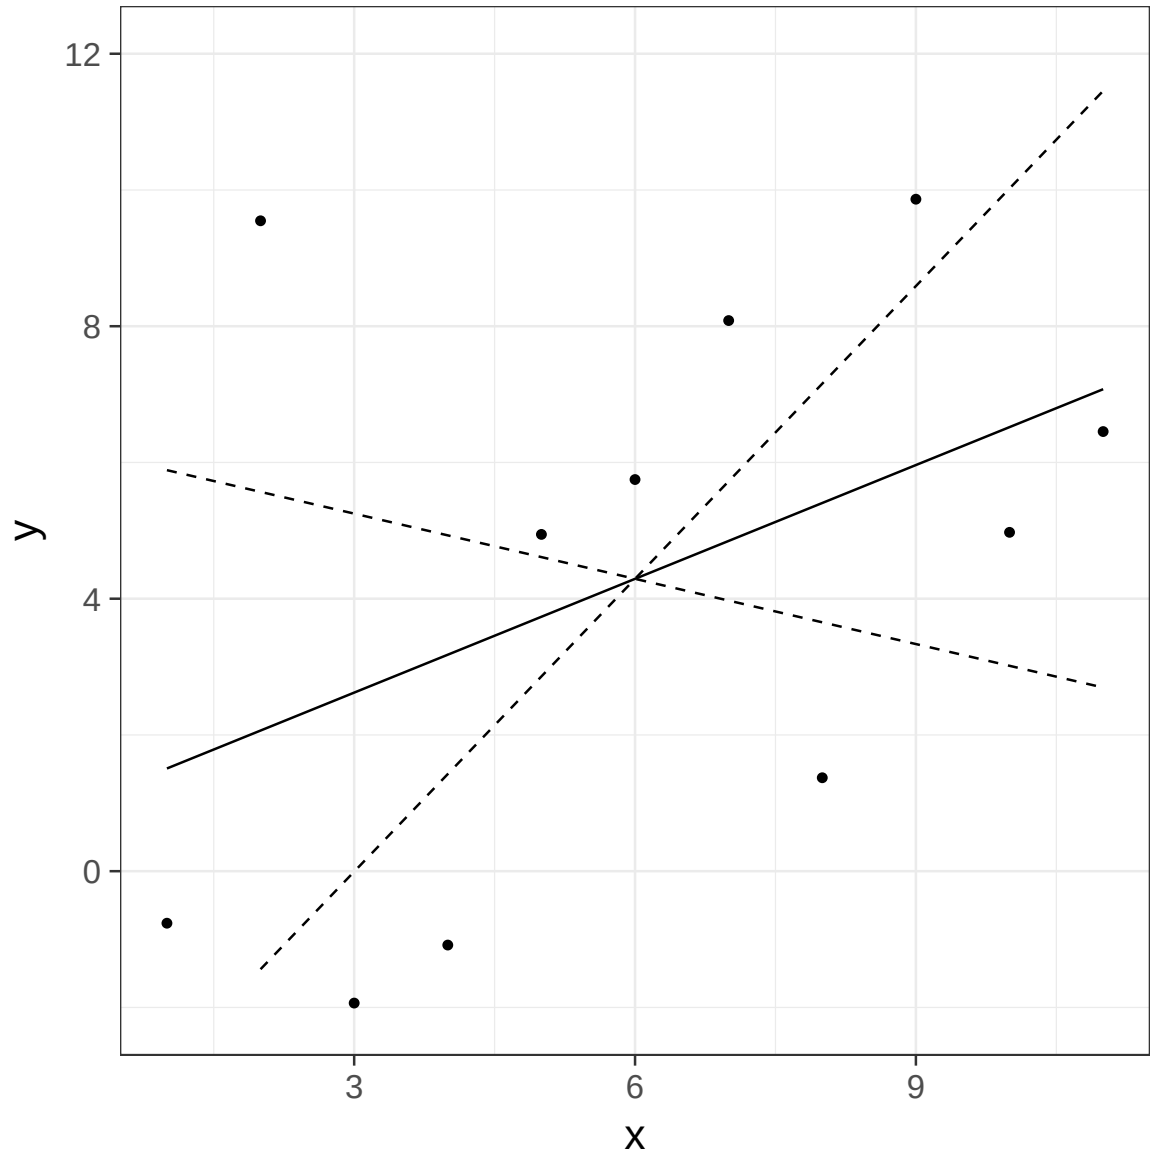

Combined

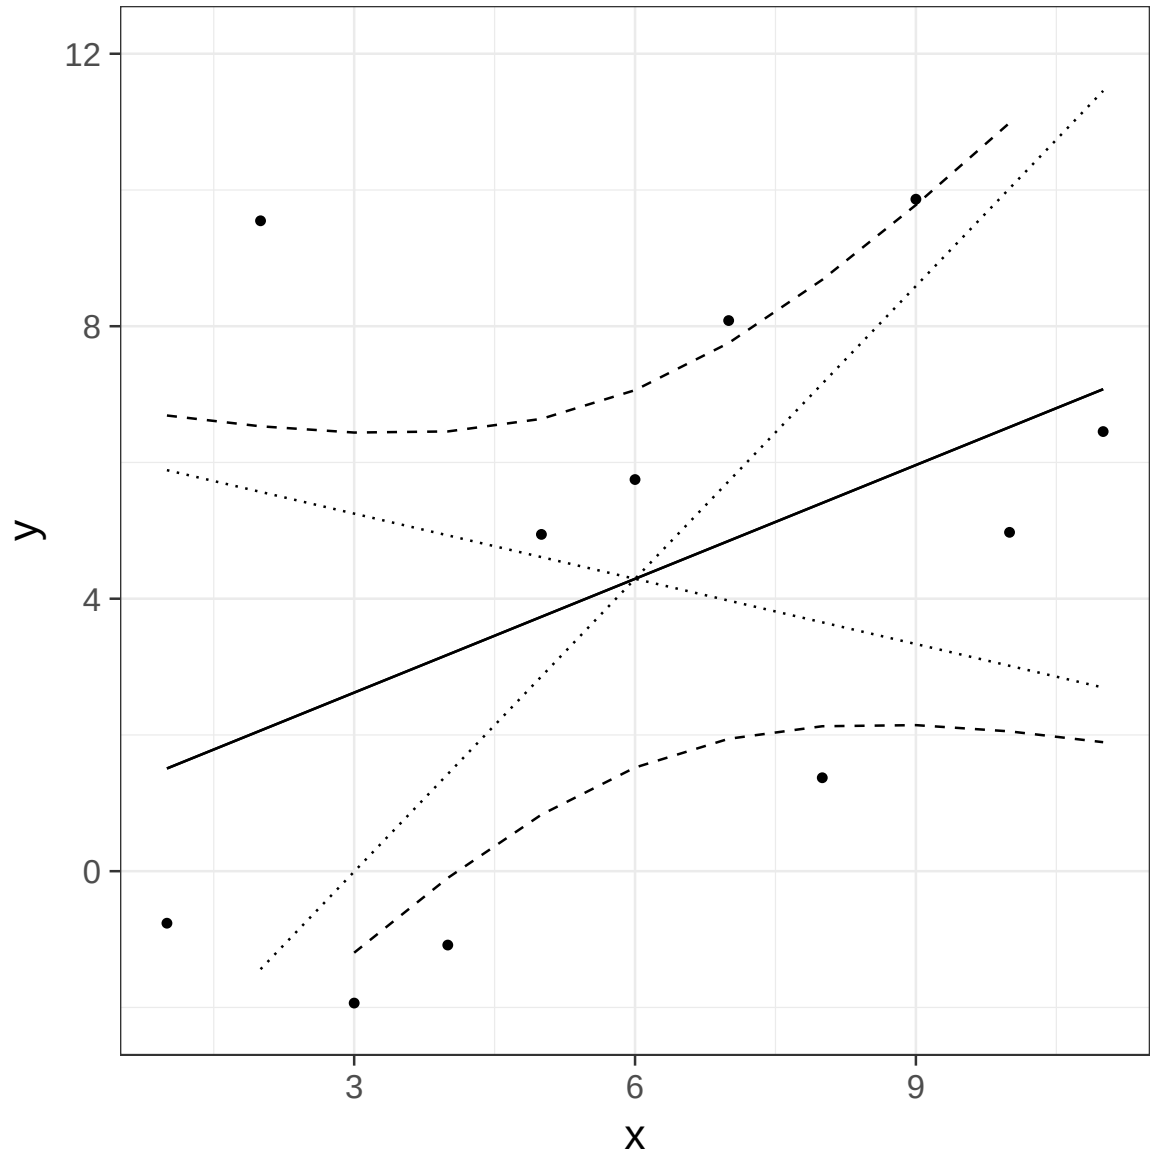

Variables vs. parameters

- A coefficient plot is most useful when each *variable* corresponds to a single statistical *parameter*
  - Binary predictor
  - Linear predictor
- More detailed shape information should be preferred when there is more than one parameter for a single logical variable
  - More than two categories
  - Splines and polynomials

No standard approach

- There are many different ways to try to capture marginal effects of a single variable
  - Particularly if it's associated with more than one parameter
- JD likes to calculate from the model "center"
  - This is the average value from each predictor column of the model matrix
  - Relatively stable
  - A bit divorced from physical reality

P values